

```
In [2]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [3]: df=pd.read_csv("C:\\Users\\asdfgh\\Desktop\\ds prodigy infotech\\twitter_data.csv")
df
```

```
Out[3]:
```

	2401	Borderlands	Positive	im getting on borderlands and i will murder you all ,
0	2401	Borderlands	Positive	I am coming to the borders and I will kill you...
1	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...
2	2401	Borderlands	Positive	im coming on borderlands and i will murder you...
3	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...
4	2401	Borderlands	Positive	im getting into borderlands and i can murder y...
...
74676	9200	Nvidia	Positive	Just realized that the Windows partition of my...
74677	9200	Nvidia	Positive	Just realized that my Mac window partition is ...
74678	9200	Nvidia	Positive	Just realized the windows partition of my Mac ...
74679	9200	Nvidia	Positive	Just realized between the windows partition of...
74680	9200	Nvidia	Positive	Just like the windows partition of my Mac is l...

74681 rows × 4 columns

```
In [4]: df.loc[-1] = df.columns
df.index = df.index + 1
df = df.sort_index()
df.columns = range(df.shape[1])
print(df)
```

	0	1	2	\
0	2401	Borderlands	Positive	
1	2401	Borderlands	Positive	
2	2401	Borderlands	Positive	
3	2401	Borderlands	Positive	
4	2401	Borderlands	Positive	
...	
74677	9200	Nvidia	Positive	
74678	9200	Nvidia	Positive	
74679	9200	Nvidia	Positive	
74680	9200	Nvidia	Positive	
74681	9200	Nvidia	Positive	

	3
0	im getting on borderlands and i will murder yo...
1	I am coming to the borders and I will kill you...
2	im getting on borderlands and i will kill you ...
3	im coming on borderlands and i will murder you...
4	im getting on borderlands 2 and i will murder ...
...	...
74677	Just realized that the Windows partition of my...
74678	Just realized that my Mac window partition is ...
74679	Just realized the windows partition of my Mac ...
74680	Just realized between the windows partition of...
74681	Just like the windows partition of my Mac is l...

[74682 rows x 4 columns]

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 74682 entries, 0 to 74681
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0    0      74682 non-null    object
1    1      74682 non-null    object
2    2      74682 non-null    object
3    3      73996 non-null    object
dtypes: object(4)
memory usage: 2.8+ MB
```

In [6]: `df.describe()`

Out[6]:

	0	1	2	3
count	74682	74682	74682	73996
unique	12448	32	4	69489
top	5203	TomClancysRainbowSix	Negative	
freq	6	2400	22542	172

In [7]: `#ID`
`df[0].unique()`

```
Out[7]: array(['2401', 2401, 2402, ..., 9198, 9199, 9200], dtype=object)
```

```
In [8]: #Entities
df[1].unique()
```

```
Out[8]: array(['Borderlands', 'CallOfDutyBlackopsColdWar', 'Amazon', 'Overwatch',
              'Xbox(Xseries)', 'NBA2K', 'Dota2', 'PlayStation5(PS5)',
              'WorldOfCraft', 'CS-GO', 'Google', 'AssassinsCreed', 'ApexLegends',
              'LeagueOfLegends', 'Fortnite', 'Microsoft', 'Hearthstone',
              'Battlefield', 'PlayerUnknownsBattlegrounds(PUBG)', 'Verizon',
              'HomeDepot', 'FIFA', 'RedDeadRedemption(RDR)', 'CallOfDuty',
              'TomClancysRainbowSix', 'Facebook', 'GrandTheftAuto(GTA)',
              'MaddenNFL', 'johnson&johnson', 'Cyberpunk2077',
              'TomClancysGhostRecon', 'Nvidia'], dtype=object)
```

```
In [9]: #Sentiments
df[2].unique()
```

```
Out[9]: array(['Positive', 'Neutral', 'Negative', 'Irrelevant'], dtype=object)
```

```
In [10]: #Tweets
df[3].unique()
```

```
Out[10]: array(['im getting on borderlands and i will murder you all ',
               'I am coming to the borders and I will kill you all,',
               'im getting on borderlands and i will kill you all,', ...,
               'Just realized the windows partition of my Mac is now 6 years behind on Nvi
               dia drivers and I have no idea how he didn't notice',
               'Just realized between the windows partition of my Mac is like being 6 year
               s behind on Nvidia drivers and cars I have no fucking idea how I ever didn ' t not
               ice',
               'Just like the windows partition of my Mac is like 6 years behind on its dr
               ivers So you have no idea how I didn't notice'],
               dtype=object)
```

```
In [11]: col_names=["Id","Entities","Sentiments","Tweets"]
df.columns=col_names
df
```

Out[11]:

	Id	Entities	Sentiments	Tweets
0	2401	Borderlands	Positive	im getting on borderlands and i will murder yo...
1	2401	Borderlands	Positive	I am coming to the borders and I will kill you...
2	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...
3	2401	Borderlands	Positive	im coming on borderlands and i will murder you...
4	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...
...
74677	9200	Nvidia	Positive	Just realized that the Windows partition of my...
74678	9200	Nvidia	Positive	Just realized that my Mac window partition is ...
74679	9200	Nvidia	Positive	Just realized the windows partition of my Mac ...
74680	9200	Nvidia	Positive	Just realized between the windows partition of...
74681	9200	Nvidia	Positive	Just like the windows partition of my Mac is l...

74682 rows × 4 columns

```
In [13]: df.isnull().sum()
```

```
Out[13]: Id          0
Entities      0
Sentiments    0
Tweets       686
dtype: int64
```

```
In [14]: df.drop_duplicates(inplace=True)
df.isnull().sum()
```

```
Out[14]: Id          0
Entities      0
Sentiments    0
Tweets       326
dtype: int64
```


```
In [15]: df.reset_index(drop=True,inplace=True)
df.isnull().sum()
```

```
Out[15]: Id          0
Entities      0
Sentiments    0
Tweets       326
dtype: int64
```


```
In [16]: sentiments_counts=df["Sentiments"].value_counts()
sentiments_counts
```

```
Out[16]: Sentiments
Negative      21787
Positive      19811
Neutral       17799
Irrelevant    12584
Name: count, dtype: int64
```

```
In [17]: plt.figure(figsize=(10,6))
sns.countplot(data=df, x="Sentiments", order=df["Sentiments"].value_counts().index)
plt.title("Sentiment Distribution")
plt.xlabel("Sentiment")
plt.show()
```




```
In [18]: c = ["red", "violet", "blue", "pink"]
plt.figure(figsize=(8,8))
plt.pie(sentiments_counts, labels=sentiments_counts.index, autopct="%1.1f%%", colors
plt.show()
```



```
In [19]: entity_sentiment_counts = df.groupby(["Entities", "Sentiments"]).size().unstack().f
print(entity_sentiment_counts)
```

Sentiments Entities	Irrelevant	Negative	Neutral	Positive
Amazon	187	566	1207	304
ApexLegends	185	579	915	610
AssassinsCreed	256	366	153	1385
Battlefield	908	449	345	565
Borderlands	238	415	584	974
CS-GO	622	337	525	723
CallOfDuty	660	866	368	428
CallOfDutyBlackopsColdWar	549	546	344	822
Cyberpunk2077	462	363	458	910
Dota2	402	706	579	542
FIFA	538	1130	100	477
Facebook	674	692	773	154
Fortnite	818	676	161	532
Google	508	571	792	339
GrandTheftAuto(GTA)	746	573	303	592
Hearthstone	218	515	687	807
HomeDepot	285	875	331	735
LeagueOfLegends	298	620	802	583
MaddenNFL	86	1671	193	373
Microsoft	169	752	823	575
NBA2K	175	1455	265	411
Nvidia	86	505	851	769
Overwatch	648	607	284	690
PlayStation5(PS5)	382	425	495	894
PlayerUnknownsBattlegrounds(PUBG)	871	664	252	380
RedDeadRedemption(RDR)	204	291	779	888
TomClancysGhostRecon	23	889	781	608
TomClancysRainbowSix	92	1113	634	505
Verizon	179	1074	555	520
WorldOfCraft	210	329	1049	716
Xbox(Xseries)	712	357	406	747
johnson&johnson	193	810	1005	253

```
In [24]: c = ["red", "violet", "blue", "green"]
entity_sentiment_counts.plot(kind="bar", stacked=True, figsize=(12, 8), color=c)
plt.title("Sentiment Analysis per Entity")
plt.xlabel("Entity")
plt.ylabel("Count of Sentiments")
plt.show()
```



```
In [20]: entity_sentiment_counts = df.groupby(["Entities", "Sentiments"]).size().unstack().f
print(entity_sentiment_counts.sort_values(by="Positive"))
```


Sentiments	Irrelevant	Negative	Neutral	Positive
Entities				
Facebook	674	692	773	154
johnson&johnson	193	810	1005	253
Amazon	187	566	1207	304
Google	508	571	792	339
MaddenNFL	86	1671	193	373
PlayerUnknownsBattlegrounds(PUBG)	871	664	252	380
NBA2K	175	1455	265	411
CallOfDuty	660	866	368	428
FIFA	538	1130	100	477
TomClancysRainbowSix	92	1113	634	505
Verizon	179	1074	555	520
Fortnite	818	676	161	532
Dota2	402	706	579	542
Battlefield	908	449	345	565
Microsoft	169	752	823	575
LeagueOfLegends	298	620	802	583
GrandTheftAuto(GTA)	746	573	303	592
TomClancysGhostRecon	23	889	781	608
ApexLegends	185	579	915	610
Overwatch	648	607	284	690
WorldOfCraft	210	329	1049	716
CS-GO	622	337	525	723
HomeDepot	285	875	331	735
Xbox(Xseries)	712	357	406	747
Nvidia	86	505	851	769
Hearthstone	218	515	687	807
CallOfDutyBlackopsColdWar	549	546	344	822
RedDeadRedemption(RDR)	204	291	779	888
PlayStation5(PS5)	382	425	495	894
Cyberpunk2077	462	363	458	910
Borderlands	238	415	584	974
AssassinsCreed	256	366	153	1385

```
In [21]: import re
from textblob import TextBlob

# Function to clean tweet text
def clean_tweet(tweet):
    tweet = tweet.lower() # Convert to Lowercase
    tweet = re.sub(r'http\S+|www\S+|https\S+', '', tweet, flags=re.MULTILINE) # Remove URLs
    tweet = re.sub(r'@\w+|\#', '', tweet) # Remove mentions and hashtags
    tweet = re.sub(r'^a-zA-Z\s', '', tweet) # Remove special characters
    return tweet

# Example tweet
example_tweet = "I love the new features of the product! Check it out: https://example.com"

# Clean the tweet
cleaned_tweet = clean_tweet(example_tweet)

# Perform sentiment analysis
blob = TextBlob(cleaned_tweet)
sentiment = blob.sentiment # Returns a tuple (polarity, subjectivity)
```

```
print(f"Cleaned Tweet: {cleaned_tweet}")
print(f"Polarity: {sentiment.polarity}, Subjectivity: {sentiment.subjectivity}")
```

Cleaned Tweet: i love the new features of the product check it out excited

Polarity: 0.3371212121212121, Subjectivity: 0.6015151515151514

```
In [22]: # Function to classify overall sentiment based on polarity
def classify_sentiment(polarity):
    if polarity > 0:
        return 'Positive'
    elif polarity < 0:
        return 'Negative'
    else:
        return 'Neutral'

# Classify the overall sentiment for the cleaned tweet
overall_sentiment = classify_sentiment(sentiment.polarity)

# Display the results
print(f"Cleaned Tweet: {cleaned_tweet}")
print(f"Polarity: {sentiment.polarity}, Subjectivity: {sentiment.subjectivity}")
print(f"Overall Sentiment: {overall_sentiment}")
```

Cleaned Tweet: i love the new features of the product check it out excited

Polarity: 0.3371212121212121, Subjectivity: 0.6015151515151514

Overall Sentiment: Positive

```
In [23]: import re

def clean_tweets(text):
    # Example of removing URLs and converting to Lowercase
    text = re.sub(r'http\S+', '', text) # Remove URLs
    text = re.sub(r'^a-zA-Z\s', '', text) # Remove non-alphabetic characters
    text = text.lower() # Convert to Lowercase
    return text

print(df.columns)
def clean_tweets(text):
    if pd.isnull(text):
        return "" # or np.nan, or any default value you want
    text = re.sub(r'http\S+', '', text)
    text = re.sub(r'^a-zA-Z\s', '', text)
    return text.lower()

df['Tweets'].dropna(inplace=True)
df['Cleaned_Tweets'] = df['Tweets'].apply(clean_tweets)
print(df['Cleaned_Tweets'])
```

```

Index(['Id', 'Entities', 'Sentiments', 'Tweets'], dtype='object')
0      im getting on borderlands and i will murder yo...
1      i am coming to the borders and i will kill you...
2      im getting on borderlands and i will kill you all
3      im coming on borderlands and i will murder you...
4      im getting on borderlands  and i will murder y...
...
71976  just realized that the windows partition of my...
71977  just realized that my mac window partition is ...
71978  just realized the windows partition of my mac ...
71979  just realized between the windows partition of...
71980  just like the windows partition of my mac is l...
Name: Cleaned_Tweets, Length: 71981, dtype: object

```

```

In [24]: from textblob import TextBlob
df['Cleaned_Tweets'] = df['Cleaned_Tweets'].fillna('')
print(df.columns)
df['Polarity'] = df['Cleaned_Tweets'].apply(lambda tweet: TextBlob(tweet).sentiment
df['Subjectivity'] = df['Cleaned_Tweets'].apply(lambda tweet: TextBlob(tweet).senti
print(df['Polarity'])
print(df['Subjectivity'])

```

```

Index(['Id', 'Entities', 'Sentiments', 'Tweets', 'Cleaned_Tweets'], dtype='object')
0      0.00
1      0.00
2      0.00
3      0.00
4      0.00
...
71976  -0.40
71977  -0.40
71978  -0.40
71979  -0.05
71980  -0.40
Name: Polarity, Length: 71981, dtype: float64
0      0.00
1      0.00
2      0.00
3      0.00
4      0.00
...
71976  0.70
71977  0.70
71978  0.70
71979  0.75
71980  0.70
Name: Subjectivity, Length: 71981, dtype: float64

```

```

In [25]: # Define classify_sentiment function
def classify_sentiment(polarity):
    if polarity > 0:
        return "Positive"
    elif polarity < 0:
        return "Negative"
    else:
        return "Neutral"

```

```

# Classify Overall Sentiment
df['Overall_Sentiment'] = df['Polarity'].apply(classify_sentiment)
print(df['Overall_Sentiment'])

# Display the DataFrame
print(df[['Tweets', 'Cleaned_Tweets', 'Polarity', 'Subjectivity', 'Overall_Sentimen

```

```

0      Neutral
1      Neutral
2      Neutral
3      Neutral
4      Neutral
...
71976  Negative
71977  Negative
71978  Negative
71979  Negative
71980  Negative
Name: Overall_Sentiment, Length: 71981, dtype: object

```

	Tweets \
0	im getting on borderlands and i will murder yo...
1	I am coming to the borders and I will kill you...
2	im getting on borderlands and i will kill you ...
3	im coming on borderlands and i will murder you...
4	im getting on borderlands 2 and i will murder ...


	Cleaned_Tweets	Polarity	Subjectivity \
0	im getting on borderlands and i will murder yo...	0.0	0.0
1	i am coming to the borders and i will kill you...	0.0	0.0
2	im getting on borderlands and i will kill you all	0.0	0.0
3	im coming on borderlands and i will murder you...	0.0	0.0
4	im getting on borderlands and i will murder y...	0.0	0.0

	Overall_Sentiment
0	Neutral
1	Neutral
2	Neutral
3	Neutral
4	Neutral


```

In [27]: plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='Overall_Sentiment', order=df['Overall_Sentiment'].value_c
plt.title('Overall Sentiment Distribution')
plt.xlabel('Sentiment')
plt.ylabel('Count')
plt.show()

```



```
In [28]: # Sentiment counts by entities
entity_sentiment_counts = df.groupby(["Entities", "Overall_Sentiment"]).size().unstack()
c = ["red", "violet", "blue"]
# Plotting
entity_sentiment_counts.plot(kind="bar", stacked=True, figsize=(12, 8), color=c)
plt.title("Sentiment Analysis per Entity")
plt.xlabel("Entity")
plt.ylabel("Count of Sentiments")
plt.legend(title='Sentiment')
plt.show()
```



```
In [29]: summary = df[['Polarity', 'Subjectivity']].describe()
print(summary)
```

	Polarity	Subjectivity
count	71981.000000	71981.000000
mean	0.072764	0.439938
std	0.339364	0.312131
min	-1.000000	0.000000
25%	-0.066667	0.166667
50%	0.000000	0.475000
75%	0.255556	0.662500
max	1.000000	1.000000

```
In [30]: # Set a threshold for negative sentiment ratio (e.g., 20%)
threshold = 20 # Adjust this value as needed

# Calculate negative sentiment ratio
negative_sentiment_count = df[df['Overall_Sentiment'] == 'Negative'].shape[0]
total_tweets = df.shape[0]
negative_sentiment_ratio = (negative_sentiment_count / total_tweets) * 100

print("Total tweets:", total_tweets)
```

```

print("Negative sentiment ratio:", negative_sentiment_ratio)

# Check if negative sentiment ratio exceeds threshold
if negative_sentiment_ratio > threshold:
    print("Alert: High negative sentiment detected!")
else:
    print("Negative sentiment is within normal range.")

```

Total tweets: 71981

Negative sentiment ratio: 28.551978994456874

Alert: High negative sentiment detected!

```

In [31]: # Display column names to ensure 'Cleaned_Tweets' exists
print("Columns:", df.columns)

# Check a sample of 'Cleaned_Tweets' to verify contents
print(df['Cleaned_Tweets'].head())

```

Columns: Index(['Id', 'Entities', 'Sentiments', 'Tweets', 'Cleaned_Tweets', 'Polarity',

```

                'Subjectivity', 'Overall_Sentiment'],
                dtype='object')
0    im getting on borderlands and i will murder yo...
1    i am coming to the borders and i will kill you...
2    im getting on borderlands and i will kill you all
3    im coming on borderlands and i will murder you...
4    im getting on borderlands and i will murder y...
Name: Cleaned_Tweets, dtype: object

```

```

In [33]: from gensim.models import LdaModel

try:
    lda_model = LdaModel(corpus, num_topics=5, id2word=dictionary, passes=15)
    topics = lda_model.print_topics(num_words=5)
    for topic in topics:
        print(topic)
except Exception as e:
    print("Error building LDA model:", e)

```

Error building LDA model: name 'corpus' is not defined

```

In [34]: import tweepy

# Replace with your own credentials
consumer_key = 'YOUR_API_KEY'
consumer_secret = 'YOUR_API_SECRET_KEY'
access_token = 'YOUR_ACCESS_TOKEN'
access_token_secret = 'YOUR_ACCESS_TOKEN_SECRET'

auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
auth.set_access_token(access_token, access_token_secret)

api = tweepy.API(auth)

# Check if you're connected
try:
    api.verify_credentials()

```

```
    print("Authentication successful")
except:
    print("Authentication failed")
```

Authentication failed

```
In [36]: from textblob import TextBlob

# Function to perform sentiment analysis
def get_sentiment(text):
    blob = TextBlob(text)
    polarity = blob.sentiment.polarity # Range: -1 (negative) to 1 (positive)
    subjectivity = blob.sentiment.subjectivity # Range: 0 (objective) to 1 (subjective)

    if polarity > 0:
        return "Positive"
    elif polarity < 0:
        return "Negative"
    else:
        return "Neutral"

# Example tweet sentiment analysis
tweet = "I love this product!"
print(f"Sentiment: {get_sentiment(tweet)}")
```

Sentiment: Positive

```
In [37]: import pandas as pd

# Sample DataFrame
df = pd.DataFrame({
    'Tweets': [
        "I love the new features of the product! Check it out: https://example.com",
        "Not satisfied with the recent update :((",
        "The product is amazing! #happy",
        "I don't like the new changes... #disappointed"
    ]
})
```

```
In [38]: print(df.columns)
```

Index(['Tweets'], dtype='object')

```
In [39]: import re
from textblob import TextBlob

# Function to clean tweet text
def clean_tweet(tweet):
    tweet = tweet.lower() # Convert to Lowercase
    tweet = re.sub(r'http\S+|www\S+|https\S+', '', tweet) # Remove URLs
    tweet = re.sub(r'@\w+|\#', '', tweet) # Remove mentions and hashtags
    tweet = re.sub(r'^a-zA-Z\s', '', tweet) # Remove special characters
    return tweet

# Clean the "Tweets" column
df['Cleaned_Tweets'] = df['Tweets'].apply(clean_tweet)
```



```
# Add polarity and subjectivity columns
df['Polarity'] = df['Cleaned_Tweets'].apply(lambda tweet: TextBlob(tweet).sentiment)
df['Subjectivity'] = df['Cleaned_Tweets'].apply(lambda tweet: TextBlob(tweet).subjectivity)

# Display the first few rows to confirm the new columns
print(df[['Tweets', 'Cleaned_Tweets', 'Polarity', 'Subjectivity']])
```

```

                                Tweets \
0  I love the new features of the product! Check ...
1          Not satisfied with the recent update :(
2          The product is amazing! #happy
3  I don't like the new changes... #disappointed

                                Cleaned_Tweets  Polarity  Subjectivity
0  i love the new features of the product check i...  0.337121      0.601515
1          not satisfied with the recent update  -0.125000      0.625000
2          the product is amazing happy  0.700000      0.950000
3  i dont like the new changes disappointed -0.306818      0.602273
```

```
In [40]: from gensim.corpora import Dictionary
# Tokenize the tweets into words
tokenized_tweets = [tweet.split() for tweet in df['Cleaned_Tweets']]
# Create a dictionary
dictionary = Dictionary(tokenized_tweets)
# Create a corpus
corpus = [dictionary.doc2bow(tweet) for tweet in tokenized_tweets]

print(tokenized_tweets)
print(dictionary)
print(corpus)
```

```
[[['i', 'love', 'the', 'new', 'features', 'of', 'the', 'product', 'check', 'it', 'out', 'excited'], ['not', 'satisfied', 'with', 'the', 'recent', 'update'], ['the', 'product', 'is', 'amazing', 'happy'], ['i', 'dont', 'like', 'the', 'new', 'changes', 'disappointed']]
Dictionary<23 unique tokens: ['check', 'excited', 'features', 'i', 'it']...>
[[ (0, 1), (1, 1), (2, 1), (3, 1), (4, 1), (5, 1), (6, 1), (7, 1), (8, 1), (9, 1), (10, 2)], [(10, 1), (11, 1), (12, 1), (13, 1), (14, 1), (15, 1)], [(9, 1), (10, 1), (16, 1), (17, 1), (18, 1)], [(3, 1), (6, 1), (10, 1), (19, 1), (20, 1), (21, 1), (22, 1)]]
```

```
In [41]: lda_model = LdaModel(corpus, num_topics=5, id2word=dictionary, passes=15)
topics = lda_model.print_topics(num_words=5)
for topic in topics:
    print(topic)
```

```
(0, '0.043*"product" + 0.043*"the" + 0.043*"amazing" + 0.043*"happy" + 0.043*"new"')
(1, '0.043*"product" + 0.043*"the" + 0.043*"is" + 0.043*"i" + 0.043*"new"')
(2, '0.043*"product" + 0.043*"the" + 0.043*"amazing" + 0.043*"is" + 0.043*"i"')
(3, '0.133*"the" + 0.072*"product" + 0.072*"i" + 0.072*"new" + 0.072*"check"')
(4, '0.142*"the" + 0.053*"dont" + 0.053*"changes" + 0.053*"like" + 0.053*"disappointed"')
```

```
In [42]: from gensim import corpora
from gensim.models import LdaModel

# Preprocess tweets
```

```

def preprocess_tweets(tweets):
    stopwords = set(["the", "and", "in", "to", "for", "of"]) # Add more stopwords
    return [[word for word in tweet.split() if word not in stopwords] for tweet in

processed_tweets = preprocess_tweets(df['Cleaned_Tweets'])

# Create dictionary and corpus
dictionary = corpora.Dictionary(processed_tweets)
corpus = [dictionary.doc2bow(tweet) for tweet in processed_tweets]

# Train LDA model
lda_model = LdaModel(corpus, num_topics=3, id2word=dictionary, passes=15)

# Display topics
topics = lda_model.print_topics(num_words=5)
for topic in topics:
    print(topic)

```

```

(0, '0.111*"with" + 0.111*"recent" + 0.111*"update" + 0.111*"not" + 0.111*"satisfie
d"')
(1, '0.117*"product" + 0.067*"features" + 0.067*"love" + 0.067*"check" + 0.067*"exci
ted"')
(2, '0.103*"i" + 0.103*"new" + 0.102*"changes" + 0.102*"disappointed" + 0.102*"lik
e"')

```

```

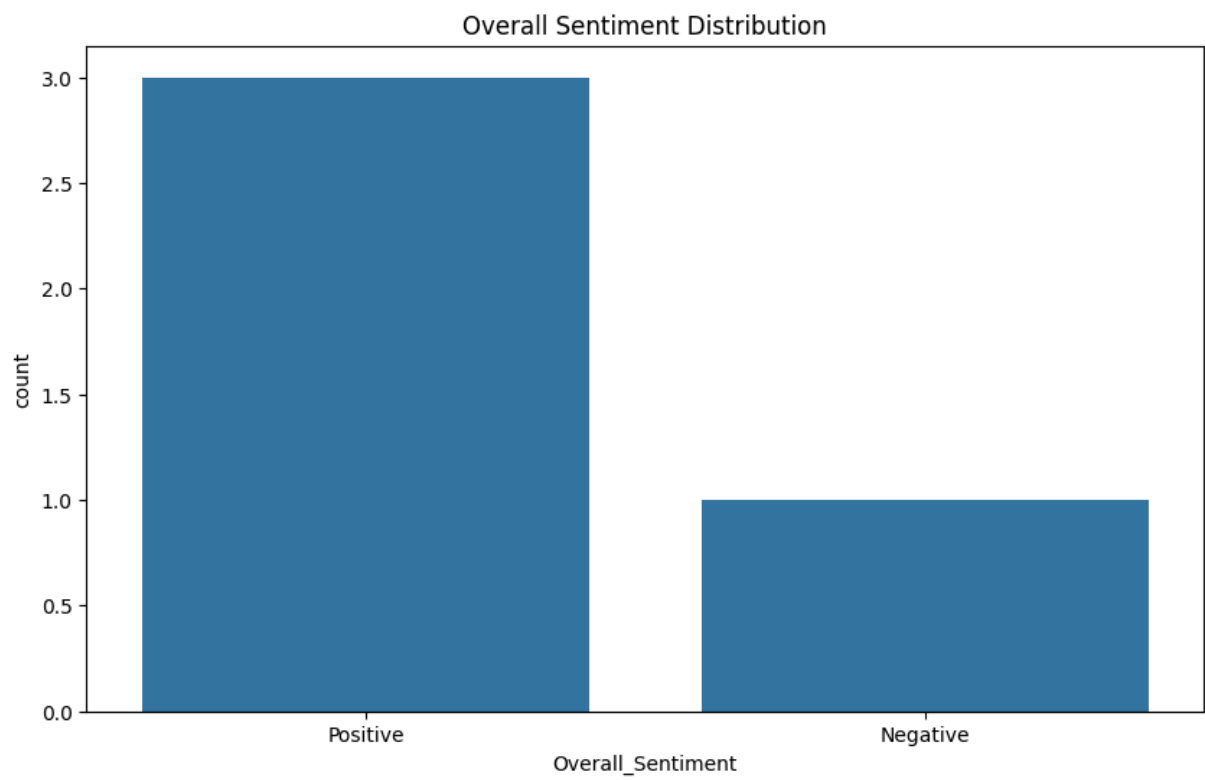
In [44]: from transformers import pipeline
# Data Preprocessing Function
def clean_text(text):
    text = text.lower()
    text = re.sub(r'http\S+|www\S+', '', text) # Remove URLs
    text = re.sub(r'@\w+|\#\w+', '', text) # Remove mentions and hashtags
    text = re.sub(r'^a-z\s', '', text) # Remove special characters
    return text

df['Cleaned_Tweets'] = df['Tweets'].apply(clean_text)

# Sentiment Analysis using TextBlob
df['Polarity'] = df['Cleaned_Tweets'].apply(lambda tweet: TextBlob(tweet).sentiment
df['Subjectivity'] = df['Cleaned_Tweets'].apply(lambda tweet: TextBlob(tweet).senti
df['Overall_Sentiment'] = df['Polarity'].apply(lambda x: 'Positive' if x > 0 else (

# Visualization
plt.figure(figsize=(10,6))
sns.countplot(data=df, x="Overall_Sentiment")
plt.title("Overall Sentiment Distribution")
plt.show()

```



In []: